

Entry-Level Machine Learning Engineer Portfolio Projects

Core Machine Learning Projects

- Movie Recommendation System
 - Skills: Collaborative filtering, matrix factorization, data wrangling
 - Stack: Python, Pandas, Scikit-learn, Surprise or LightFM
 - Bonus: Deploy with Streamlit or Flask
- Used Car Price Prediction
 - Skills: Regression, feature engineering, outlier handling
 - Stack: Python, Scikit-learn, Pandas, NumPy
 - Bonus: Host as an API and serve predictions
- Fake News Detection
 - Skills: NLP, text preprocessing, TF-IDF, logistic regression
 - Stack: Python, Scikit-learn, NLTK or SpaCy
 - Bonus: Add explainability with LIME or SHAP
- Customer Segmentation
 - Skills: Clustering (KMeans, DBSCAN), PCA
 - Stack: Python, Pandas, Matplotlib, Seaborn
 - Bonus: Use real-world datasets like Mall Customers or E-Commerce logs
- Handwritten Digit Recognition (MNIST)
 - Skills: Neural networks, CNNs
 - Stack: TensorFlow or PyTorch
 - Bonus: Create a simple web app that classifies drawn digits

Real-World Simulation Projects

- Loan Eligibility Prediction

- Skills: Classification, dealing with imbalanced data (SMOTE), model evaluation
- Stack: Python, Scikit-learn, XGBoost
- Bonus: Add explainable ML to show why a loan was rejected
- Time Series Forecasting Stock or Energy Prices
 - Skills: ARIMA, LSTM, feature windowing
 - Stack: Python, statsmodels, Keras/TensorFlow
 - Bonus: Visualize predictions with Plotly
- Real-time Object Detection
 - Skills: Computer vision, YOLO or MobileNet
 - Stack: Python, OpenCV, TensorFlow Lite
 - Bonus: Run on webcam or Raspberry Pi

Capstone-Level Portfolio Projects

- AI Resume Screener
 - Skills: NLP + classification (good/bad resume), keyword extraction
 - Stack: Python, HuggingFace Transformers
 - Bonus: Use Streamlit to create a frontend
- ML-Powered Personal Finance Tracker
 - Skills: Text classification (labeling expenses), clustering
 - Stack: Python, FastAPI, SQLite
 - Bonus: Create a full-stack version with a simple dashboard

Tips to Stand Out

Use GitHub: Include clean code, README, and visualizations.

Include Metrics: Always report accuracy, F1, confusion matrix, etc.

Deploy at least one project: Use Streamlit, HuggingFace Spaces, or Render.

Write a blog post for at least 12 projects to explain your thought process.